



↳ 1 cells hidden

Creating and Manipulating Tensors

Learning Objectives:

- Initialize and assign TensorFlow Variables
- Create and manipulate tensors
- Refresh your memory about addition and multiplication in linear algebra (consult an introduction to matrix [addition](#) and [multiplication](#) if these topics are new to you)
- Familiarize yourself with basic TensorFlow math and array operations

```
import tensorflow as tf
```

Vector Addition

You can perform many typical mathematical operations on tensors ([TF API](#)). The following code creates and manipulates two vectors (1-D tensors), each having exactly six elements:

```
with tf.Graph().as_default():
    # Create a six-element vector (1-D tensor).
    primes = tf.constant([2, 3, 5, 7, 11, 13], dtype=tf.int32)

    # Create another six-element vector. Each element in the vector will be
    # initialized to 1. The first argument is the shape of the tensor (more
    # on shapes below).
    ones = tf.ones([6], dtype=tf.int32)

    # Add the two vectors. The resulting tensor is a six-element vector.
    just_beyond_primes = tf.add(primes, ones)

    # Create a session to run the default graph.
    with tf.Session() as sess:
        print just_beyond_primes.eval()
```

Tensor Shapes

Shapes are used to characterize the size and number of dimensions of a tensor. The shape of a tensor is expressed as a list, with the *i*th element representing the size along dimension *i*. The length of the list then indicates the rank of the tensor (i.e., the number of dimensions).

For more information, see the [TensorFlow documentation](#).

A few basic examples:

```
with tf.Graph().as_default():
    # A scalar (0-D tensor).
    scalar = tf.zeros([])

    # A vector with 3 elements.
    vector = tf.zeros([3])

    # A matrix with 2 rows and 3 columns.
    matrix = tf.zeros([2, 3])

    with tf.Session() as sess:
        print 'scalar has shape', scalar.get_shape(), 'and value:\n', scalar.eval()
```

```
print 'vector has shape', vector.get_shape(), 'and value:\n', vector.eval()  
print 'matrix has shape', matrix.get_shape(), 'and value:\n', matrix.eval()
```

▼ Broadcasting

In mathematics, you can only perform element-wise operations (e.g. *add* and *equals*) on tensors of the same shape. In TensorFlow, however, you may perform operations on tensors that would traditionally have been incompatible.

TensorFlow supports **broadcasting** (a concept borrowed from numpy), where the smaller array in an element-wise operation is enlarged to have the same shape as the larger array. For example, via broadcasting:

- If an operand requires a size [6] tensor, a size [1] or a size [] tensor can serve as an operand.
- If an operation requires a size [4, 6] tensor, any of the following sizes can serve as an operand:
 - [1, 6]
 - [6]
 - []
- If an operation requires a size [3, 5, 6] tensor, any of the following sizes can serve as an operand:
 - [1, 5, 6]
 - [3, 1, 6]
 - [3, 5, 1]
 - [1, 1, 1]
 - [5, 6]
 - [1, 6]
 - [6]
 - [1]
 - []

NOTE: When a tensor is broadcast, its entries are conceptually **copied**. (They are not actually copied for performance reasons. Broadcasting was invented as a performance optimization.)

The full broadcasting ruleset is well described in the easy-to-read [numpy broadcasting documentation](#).

The following code performs the same tensor addition as before, but using broadcasting:

```
with tf.Graph().as_default():  
    # Create a six-element vector (1-D tensor).  
    primes = tf.constant([2, 3, 5, 7, 11, 13], dtype=tf.int32)  
  
    # Create a constant scalar with value 1.  
    ones = tf.constant(1, dtype=tf.int32)  
  
    # Add the two tensors. The resulting tensor is a six-element vector.  
    just_beyond_primes = tf.add(primes, ones)  
  
    with tf.Session() as sess:  
        print just_beyond_primes.eval()
```

▼ Matrix Multiplication

In linear algebra, when multiplying two matrices, the number of *columns* of the first matrix must equal the number of *rows* in the second matrix.

- It is **valid** to multiply a 3x4 matrix by a 4x2 matrix. This will result in a 3x2 matrix.
- It is **invalid** to multiply a 4x2 matrix by a 3x4 matrix.

```
with tf.Graph().as_default():  
    # Create a matrix (2-d tensor) with 3 rows and 4 columns.  
    x = tf.constant([[5, 2, 4, 3], [5, 1, 6, -2], [-1, 3, -1, -2]],  
                    dtype=tf.int32)  
  
    # Create a matrix with 4 rows and 2 columns.
```

```
y = tf.constant([[2, 2], [3, 5], [4, 5], [1, 6]], dtype=tf.int32)
```

```
# Multiply `x` by `y`.
# The resulting matrix will have 3 rows and 2 columns.
matrix_multiply_result = tf.matmul(x, y)
```

```
with tf.Session() as sess:
    print matrix_multiply_result.eval()
```

▼ Tensor Reshaping

With tensor addition and matrix multiplication each imposing constraints on operands, TensorFlow programmers must frequently reshape tensors.

You can use the `tf.reshape` method to reshape a tensor. For example, you can reshape a 8x2 tensor into a 2x8 tensor or a 4x4 tensor:

```
with tf.Graph().as_default():
    # Create an 8x2 matrix (2-D tensor).
    matrix = tf.constant([[1,2], [3,4], [5,6], [7,8],
                          [9,10], [11,12], [13, 14], [15,16]], dtype=tf.int32)

    # Reshape the 8x2 matrix into a 2x8 matrix.
    reshaped_2x8_matrix = tf.reshape(matrix, [2,8])

    # Reshape the 8x2 matrix into a 4x4 matrix
    reshaped_4x4_matrix = tf.reshape(matrix, [4,4])

    with tf.Session() as sess:
        print "Original matrix (8x2):"
        print matrix.eval()
        print "Reshaped matrix (2x8):"
        print reshaped_2x8_matrix.eval()
        print "Reshaped matrix (4x4):"
        print reshaped_4x4_matrix.eval()
```

You can also use `tf.reshape` to change the number of dimensions (the "rank") of the tensor. For example, you could reshape that 8x2 tensor into a 3-D 2x2x4 tensor or a 1-D 16-element tensor.

```
with tf.Graph().as_default():
    # Create an 8x2 matrix (2-D tensor).
    matrix = tf.constant([[1,2], [3,4], [5,6], [7,8],
                          [9,10], [11,12], [13, 14], [15,16]], dtype=tf.int32)

    # Reshape the 8x2 matrix into a 3-D 2x2x4 tensor.
    reshaped_2x2x4_tensor = tf.reshape(matrix, [2,2,4])

    # Reshape the 8x2 matrix into a 1-D 16-element tensor.
    one_dimensional_vector = tf.reshape(matrix, [16])

    with tf.Session() as sess:
        print "Original matrix (8x2):"
        print matrix.eval()
        print "Reshaped 3-D tensor (2x2x4):"
        print reshaped_2x2x4_tensor.eval()
        print "1-D vector:"
        print one_dimensional_vector.eval()
```

▼ Exercise #1: Reshape two tensors in order to multiply them.

The following two vectors are incompatible for matrix multiplication:

- `a = tf.constant([5, 3, 2, 7, 1, 4])`
- `b = tf.constant([4, 6, 3])`

Reshape these vectors into compatible operands for matrix multiplication. Then, invoke a matrix multiplication operation on the reshaped tensors.

```
# Write your code for Task 1 here.
```

▼ Solution

Click below for a solution.

```
with tf.Graph().as_default(), tf.Session() as sess:
    # Task: Reshape two tensors in order to multiply them

    # Here are the original operands, which are incompatible
    # for matrix multiplication:
    a = tf.constant([5, 3, 2, 7, 1, 4])
    b = tf.constant([4, 6, 3])
    # We need to reshape at least one of these operands so that
    # the number of columns in the first operand equals the number
    # of rows in the second operand.

    # Reshape vector "a" into a 2-D 2x3 matrix:
    reshaped_a = tf.reshape(a, [2,3])

    # Reshape vector "b" into a 2-D 3x1 matrix:
    reshaped_b = tf.reshape(b, [3,1])

    # The number of columns in the first matrix now equals
    # the number of rows in the second matrix. Therefore, you
    # can matrix multiply the two operands.
    c = tf.matmul(reshaped_a, reshaped_b)
    print(c.eval())

    # An alternate approach: [6,1] x [1, 3] -> [6,3]
```

▼ Variables, Initialization and Assignment

So far, all the operations we performed were on static values (`tf.constant`); calling `eval()` always returned the same result. TensorFlow allows you to define `Variable` objects, whose values can be changed.

When creating a variable, you can set an initial value explicitly, or you can use an initializer (like a distribution):

```
g = tf.Graph()
with g.as_default():
    # Create a variable with the initial value 3.
    v = tf.Variable([3])

    # Create a variable of shape [1], with a random initial value,
    # sampled from a normal distribution with mean 1 and standard deviation 0.35.
    w = tf.Variable(tf.random_normal([1], mean=1.0, stddev=0.35))
```

One peculiarity of TensorFlow is that **variable initialization is not automatic**. For example, the following block will cause an error:

```
with g.as_default():
    with tf.Session() as sess:
        try:
            v.eval()
        except tf.errors.FailedPreconditionError as e:
            print "Caught expected error: ", e
```

The easiest way to initialize a variable is to call `global_variables_initializer`. Note the use of `Session.run()`, which is roughly equivalent to `eval()`.

```
with g.as_default():
    with tf.Session() as sess:
        initialization = tf.global_variables_initializer()
        sess.run(initialization)
        # Now, variables can be accessed normally, and have values assigned to them.
```

```
print v.eval()  
print w.eval()
```

Once initialized, variables will maintain their value within the same session (however, when starting a new session, you will need to re-initialize them):

```
with g.as_default():  
    with tf.Session() as sess:  
        sess.run(tf.global_variables_initializer())  
        # These three prints will print the same value.  
        print w.eval()  
        print w.eval()  
        print w.eval()
```

To change the value of a variable, use the assign op. Note that simply creating the assign op will not have any effect. As with initialization, you have to run the assignment op to update the variable value:

```
with g.as_default():  
    with tf.Session() as sess:  
        sess.run(tf.global_variables_initializer())  
        # This should print the variable's initial value.  
        print v.eval()  
  
        assignment = tf.assign(v, [7])  
        # The variable has not been changed yet!  
        print v.eval()  
  
        # Execute the assignment op.  
        sess.run(assignment)  
        # Now the variable is updated.  
        print v.eval()
```

There are many more topics about variables that we didn't cover here, such as loading and storing. To learn more, see the [TensorFlow docs](#).

▼ Exercise #2: Simulate 10 rolls of two dice.

Create a dice simulation, which generates a 10x3 2-D tensor in which:

- Columns 1 and 2 each hold one throw of one six-sided die (with values 1–6).
- Column 3 holds the sum of Columns 1 and 2 on the same row.

For example, the first row might have the following values:

- Column 1 holds 4
- Column 2 holds 3
- Column 3 holds 7

You'll need to explore the [TensorFlow documentation](#) to solve this task.

```
# Write your code for Task 2 here.
```

▼ Solution

Click below for a solution.

```
with tf.Graph().as_default(), tf.Session() as sess:  
    # Task 2: Simulate 10 throws of two six-sided dice. Store the results  
    # in a 10x3 matrix.  
  
    # We're going to place dice throws inside two separate  
    # 10x1 matrices. We could have placed dice throws inside  
    # a single 10x2 matrix, but adding different columns of
```

```
# the same matrix is tricky. We also could have placed
# dice throws inside two 1-D tensors (vectors); doing so
# would require transposing the result.
dice1 = tf.Variable(tf.random_uniform([10, 1],
                                     minval=1, maxval=7,
                                     dtype=tf.int32))
dice2 = tf.Variable(tf.random_uniform([10, 1],
                                     minval=1, maxval=7,
                                     dtype=tf.int32))

# We may add dice1 and dice2 since they share the same shape
# and size.
dice_sum = tf.add(dice1, dice2)

# We've got three separate 10x1 matrices. To produce a single
# 10x3 matrix, we'll concatenate them along dimension 1.
resulting_matrix = tf.concat(
    values=[dice1, dice2, dice_sum], axis=1)

# The variables haven't been initialized within the graph yet,
# so let's remedy that.
sess.run(tf.global_variables_initializer())

print(resulting_matrix.eval())
```